


# UAV-Net: A Fast Aerial Vehicle Detector for Mobile Platforms

## 3rd International Workshop on Computer Vision for UAVs – CVPR 2019

**Tobias Ringwald, Lars Sommer, Arne Schumann, Jürgen Beyerer and Rainer Stiefelhagen**

Karlsruhe Institute of Technology – Institute for Anthropomatics and Robotics, Computer Vision for Human-Computer Interaction  
Fraunhofer Institute of Optronics, System Technologies and Image Exploitation




# Motivation

- Deep learning best solution for object detection
- Large server clusters for training and inference
- „Intelligence“ also desired in edge devices
- Problems with weight, power supply and dimensions



+




=

?

# Solution

- Jetson platform by NVIDIA
- For use in „intelligent“ cars, cameras, **drones** etc.
- Embedded GPU with cuDNN stack
- Jetson TX2:
  - 8GB RAM
  - 6-core CPU @ 2GHz
  - 256 CUDA cores
  - Max. 15W
- Is it enough?






# UAV-Net

- Small and efficient detector for on-board object detection
- Very low memory footprint
- On par with state-of-the-art detection models
- Evaluated on 3 different datasets
- Design decisions: **Meta-architecture, backbone, layers, filters**



DLR 3K Munich



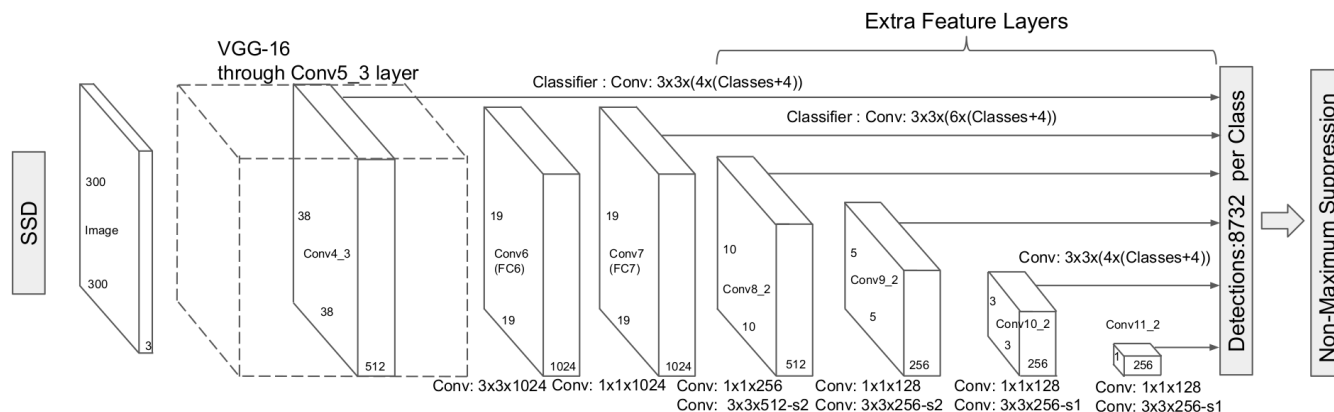
VEDAI



UAVDT


# Meta-architectures

- Candidates: Faster R-CNN<sup>[1]</sup>, SSD<sup>[2]</sup> and YOLOv2<sup>[3]</sup>
- SSD offers best trade-off, YOLOv2 competitive
- Quick SSD recap:
  - Base network (backbone), initially VGG-16
  - Extra feature layers
  - Convolutional layers for classification and box regression
  - Non-maximum suppression




# Meta-Architecture Modifications

- SSD makes use of multiple scales
  - Constant GSD
  - Use 8× downsampled feature maps
  - Found by ablation study
- Box sizes and ratios
  - Boxes have to fit the object size
  - Clustering approach similar to DSSD<sup>[4]</sup>
  - Saves filters in the prediction layers with next to no change in accuracy




# Base Networks

- Many different networks in literature
  - MobileNet<sup>[5]</sup>
  - ShuffleNet<sup>[6]</sup>
  - SqueezeNet<sup>[7]</sup>
  - ZynqNet<sup>[8]</sup>
- ZynqNet (SqueezeNet-like architecture)
  - Only standard layers
  - Strided convolution instead of pooling
  - Alternating 3×3 and 1×1 for squeeze layer
- Additional changes
  - No ReLU after squeeze layer
  - ELU instead of ReLUs



# Regression and Classification Layers


- SSD uses  $3 \times 3$  convolutions by default
- But strongest feature for vehicles is the windshield
- On a  $8 \times$  downsampled feature map only 1 „pixel“ is covered
- $1 \times 1$  convolutions are sufficient
- No loss in average precission





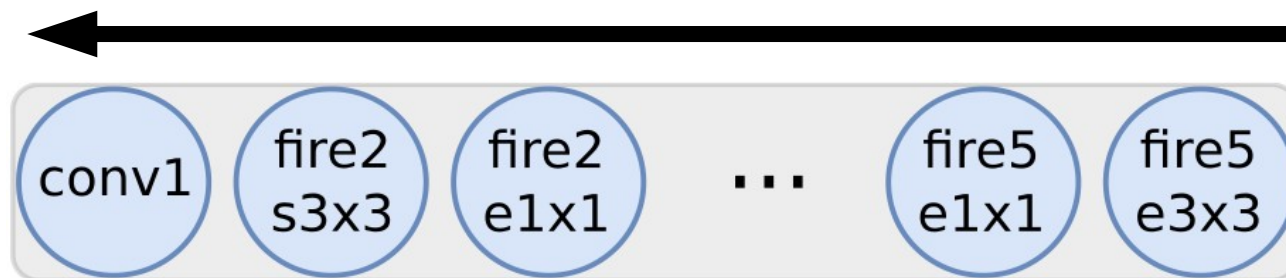
# Auto-Pruning

- Objective is to reduce number of filters in the base network from  $N$  to  $\phi \times N$




# Auto-Pruning


- Iterate from back to front (later layers have more filters)
- Delete  $k$  filters in layer  $L_n$ , according to  $l_1$  norm
- Calculate validation sensitivity  $S$
- Yielding tuples of  $(S, L_n, k)$
- Remove worst tuple from network according to metric
- No retraining required!



# Auto-Pruning




# Auto-Pruning





# Auto-Pruning



# Auto-Pruning

- Target is predefined  $\phi$  value
  - Only one hyperparameter
  - Arbitrarily chosen: 0.5, 0.25, 0.15 etc.
  - Can be adjusted for use case
- Pruning decision can be any metric, also depending on application area (speed vs. accuracy)
- Final network only has to be finetuned once for a few iterations

# Quantitative Results

- DLR 3K results:

| Network                      | AP (%)      | Inference Speed (FPS) |              |             |
|------------------------------|-------------|-----------------------|--------------|-------------|
|                              |             | Titan X               | GTX 1060     | Jetson TX2  |
| VGG                          | <b>97.2</b> | 27.9                  | 11.7         | 1.3         |
| ZynqNet                      | <b>97.2</b> | 184.9                 | 81.9         | 14.7        |
| UAV-Net $_{\varphi = 1.000}$ | <b>97.2</b> | 194.1                 | 83.8         | 15.9        |
| UAV-Net $_{\varphi = 0.750}$ | <b>97.2</b> | 225.8                 | 98.8         | 18.8        |
| UAV-Net $_{\varphi = 0.500}$ | 97.1        | 265.2                 | 116.2        | 22.7        |
| UAV-Net $_{\varphi = 0.250}$ | 95.4        | 342.6                 | 153.3        | 31.3        |
| UAV-Net $_{\varphi = 0.150}$ | 91.3        | 410.0                 | 181.2        | 38.2        |
| UAV-Net $_{\varphi = 0.075}$ | 11.1        | <b>426.8</b>          | <b>203.5</b> | <b>43.1</b> |

# Quantitative Results

- DLR 3K results:

| Network                      | AP (%)      | Inference Speed (FPS) |              |             |
|------------------------------|-------------|-----------------------|--------------|-------------|
|                              |             | Titan X               | GTX 1060     | Jetson TX2  |
| VGG                          | <b>97.2</b> | 27.9                  | 11.7         | 1.3         |
| ZynqNet                      | <b>97.2</b> | 184.9                 | 81.9         | 14.7        |
| UAV-Net $_{\varphi = 1.000}$ | <b>97.2</b> | 194.1                 | 83.8         | 15.9        |
| UAV-Net $_{\varphi = 0.750}$ | <b>97.2</b> | 225.8                 | 98.8         | 18.8        |
| UAV-Net $_{\varphi = 0.500}$ | 97.1        | 265.2                 | 116.2        | 22.7        |
| UAV-Net $_{\varphi = 0.250}$ | 95.4        | 342.6                 | 153.3        | 31.3        |
| UAV-Net $_{\varphi = 0.150}$ | 91.3        | 410.0                 | 181.2        | 38.2        |
| UAV-Net $_{\varphi = 0.075}$ | 11.1        | <b>426.8</b>          | <b>203.5</b> | <b>43.1</b> |

| Network architecture        | Model Size      | Parameter Count | Relative Size |
|-----------------------------|-----------------|-----------------|---------------|
| VGG, 2 box sizes            | 30.19 MiB       | 7,912,316       | 100.0%        |
| ZynqNet                     | 0.89 MiB        | 230,782         | 2.9%          |
| UAV-Net $_{\varphi = 0.50}$ | 0.39 MiB        | 101,934         | 1.3%          |
| UAV-Net $_{\varphi = 0.15}$ | <b>0.07 MiB</b> | <b>17,146</b>   | <b>0.2%</b>   |



# Quantitative Results

- VEDAI-1024 results:

| Dataset | Model                     | AP (%)      | Inference Speed (FPS) |              |             |
|---------|---------------------------|-------------|-----------------------|--------------|-------------|
|         |                           |             | Titan X               | GTX 1060     | Jetson TX2  |
| VEDAI   | VGG                       | <b>96.4</b> | 16.7                  | 5.8          | 0.7         |
| VEDAI   | UAV-Net $_{\varphi=1.00}$ | 95.7        | 123.5                 | 50.2         | 9.9         |
| VEDAI   | UAV-Net $_{\varphi=0.50}$ | 95.2        | 168.0                 | 73.8         | 13.9        |
| VEDAI   | UAV-Net $_{\varphi=0.15}$ | 93.5        | <b>256.4</b>          | <b>125.9</b> | <b>22.9</b> |

# What about other setups?

- Some modifications specific to dataset: constant GSD, constant object sizes etc.
- Evaluation also shown for UAVDT
  - 1x1 regression filters not large enough
  - Single box size not sufficient anymore
  - Other modifications still useful



DLR 3K Munich



UAVDT

# Quantitative Results

- UAVDT results:

| Dataset | Model                                       | AP (%)       | Inference Speed (FPS) |             |             |
|---------|---|--------------|-----------------------|-------------|-------------|
|         |   |              | Titan X               | GTX 1060    | Jetson TX2  |
| UAVDT   | R-FCN [17]                                  | 34.35        | 4.7                   | —           | —           |
| UAVDT   | SSD [17]                                    | 33.62        | 41.6                  | —           | —           |
| UAVDT   | Faster R-CNN [17]                           | 22.32        | 2.8                   | —           | —           |
| UAVDT   | RON [17]                                    | 21.59        | 11.1                  | —           | —           |
| UAVDT   | UAV-Net $_{\varphi=1.00}^{1 \times 1, c=1}$ | 26.21        | <b>214.0</b>          | <b>98.8</b> | <b>18.3</b> |
| UAVDT   | UAV-Net $_{\varphi=1.00}^{5 \times 5, c=5}$ | <b>34.52</b> | 80.1                  | 34.7        | 6.6         |
| UAVDT   | UAV-Net $_{\varphi=1.00}^{3 \times 3, c=4}$ | 32.76        | 112.2                 | 51.5        | 9.0         |
| UAVDT   | UAV-Net $_{\varphi=0.50}^{3 \times 3, c=4}$ | 31.82        | 132.5                 | 69.2        | 11.4        |

# Qualitative Results – DLR 3K



UAV-Net <sub>$\varphi=0.50$</sub>

UAV-Net <sub>$\varphi=0.15$</sub>



# Qualitative Results – UAVDT




UAV-Net <sub>$\varphi=1.00$</sub>  (5×5, c=5)

UAV-Net <sub>$\varphi=0.50$</sub>  (3×3, c=4)

# UAV-Net Summary

- Fast but still accurate vehicle detector
  - 38 FPS on DLR 3K with >90% mAP (test set)
- Ultra-low footprint for both model size and memory usage
  - As low as 0.07 MiB for DLR 3K and VEDAI
- Power envelope of <15W



# Image Sources

- Title Slide: "Artificial Intelligence & AI & Machine Learning" by mikemacmarketing is licensed under CC BY 2.0
- Drone: <https://fortunedotcom.files.wordpress.com/2015/12/drone.jpg>
- NVIDIA Titan V:  
[https://www.nvidia.com/content/dam/en-zz/es\\_em/Solutions/geforce/TITAN/TITANV/nvidia-titan-xp-shop-625-ud.jpg](https://www.nvidia.com/content/dam/en-zz/es_em/Solutions/geforce/TITAN/TITANV/nvidia-titan-xp-shop-625-ud.jpg)
- Jetson Platform:  
[https://devblogs.nvidia.com/wp-content/uploads/2017/03/Figure1\\_TX2-e1488772330657.png](https://devblogs.nvidia.com/wp-content/uploads/2017/03/Figure1_TX2-e1488772330657.png)  
[https://devblogs.nvidia.com/wp-content/uploads/2017/03/JTX2\\_Devkit-e1488775199359-624x615.png](https://devblogs.nvidia.com/wp-content/uploads/2017/03/JTX2_Devkit-e1488775199359-624x615.png)

# References

- [1] S. Ren, K. He, R. Girshick, and J. Sun. Faster R-CNN: Towards real-time object detection with region proposal networks. In NIPS, 2015
- [2] W. Liu, D. Anguelov, D. Erhan, C. Szegedy, S. E. Reed, C. Fu, and A. C. Berg, “SSD: single shot multibox detector,” CoRR, vol. abs/1512.02325, 2015.
- [3] J. Redmon and A. Farhadi. YOLO9000: better, faster, stronger. CVPR, 2017.
- [4] C.-Y. Fu, W. Liu, A. Ranga, A. Tyagi, and A. C. Berg. DSSD: Deconvolutional single shot detector. arXiv preprint arXiv:1701.06659, 2017
- [5] A. G. Howard, M. Zhu, B. Chen, D. Kalenichenko, W. Wang, T. Weyand, M. Andreetto, and H. Adam, “Mobilenets: Efficient convolutional neural networks for mobile vision applications,” CoRR, vol. abs/1704.04861, 2017.
- [6] X. Zhang, X. Zhou, M. Lin, and J. Sun, “Shufflenet: An extremely efficient convolutional neural network for mobile devices,” CoRR, vol. abs/1707.01083, 2017.
- [7] F. N. Iandola, M. W. Moskewicz, K. Ashraf, S. Han, W. J. Dally, and K. Keutzer, “Squeezenet: Alexnet-level accuracy with 50x fewer parameters and <1mb model size,” CoRR, vol. abs/1602.07360, 2016.
- [8] D. Gschwend, “Zynqnet: An fpga-accelerated embedded convolutional neural network,” Swiss Federal Institute of Technology Zurich (ETH-Zurich), 2016.



# Additional Slides


# Qualitative Results – DLR 3K



UAV-Net <sub>$\varphi=0.50$</sub>

UAV-Net <sub>$\varphi=0.15$</sub>



# Qualitative Results – DLR 3K Error Cases



UAV-Net <sub>$\varphi=0.50$</sub>



# Qualitative Results – UAVDT



UAV-Net <sub>$\varphi=1.00$</sub>  (5×5,c=5)

UAV-Net <sub>$\varphi=0.50$</sub>  (3×3,c=4)